

CSC1202, 2026 – Week 2: AI Agent Architecture and Lifecycle

- **Part 1: LLM AI Agent Architecture**
- **Part 2: LLM AI Agent lifecycle**
- **Part 3: LLM Mult-agent system**
- **References: REF1.** Grootendorst, M., & Alammari, J. (2026). An Illustrated Guide to AI Agents: Concepts and Code for Building Agents with LLMs, Tools, and Memory. O'Reilly Media. **REF2.** Selected research papers, technical reports, and practitioner resources on LLM-based agents.

Why Do We Need LLM Agents?

- LLMs alone are limited when tasks require data, actions, multiple steps and continuity.
- **Need access to external data and knowledge**
 - Enterprise tasks often require databases, CRM data, documents, or current information.
- **Need to handle multi-step tasks**
 - Real tasks require several connected steps (e.g., find flights, compare options, select, and book).
- **Need access to tools and applications**
 - Completing tasks may require APIs, databases, calendars, CRM systems, or other applications.
- **Need to maintain context and memory**
 - Information from earlier steps or interactions may be needed later.

From LLMs to LLM Agents

LLM: Generate a Response

User
Prompt

LLM

Generated
Response

The interaction can finish once the model produces the response (stateless)

LLM-Based Agent: Pursue a Goal / Perform a Task

Goal /
Task

**LLM-
Based
Agent**

**Environ
ment**

Action
Observation
e.g., APIs • Web
Apps • Databases

The agent can decide what to do next, take an action, observe the result, and continue until the task is completed.

Key distinction

An LLM generates a response. An LLM-based agent uses an LLM to decide and act while interacting with an environment to accomplish a goal.

Example: Google NotebookLM

- **Purpose:** AI research assistant that helps users understand, analyse and create content from their own sources and information discovered from the Web.
- **Model: Gemini.**
- **Context / Knowledge:** Uses PDFs, documents, slides, websites, YouTube, audio and other sources collected in a notebook.
- **Example actions:** search the Web, discover/select sources, retrieve information from sources, analyse information, and generate reports, summaries, audio overviews and other outputs.
- **How it works:** Given a research task, it can **find/use relevant sources, reason over the information, perform actions as needed, generate a grounded output with citations**


Example: Claude Code — Coding Agent

- **Purpose:** Agentic coding assistant that helps developers build, debug, test and maintain software.
- **Model:** Claude.
- **Context / Environment:** Works with the developer's **codebase, files, project instructions and development environment.**
- **Example actions/tools:** read/search files, create or modify code, run commands, execute tests, use Git, and interact with development tools.
- **How it works:** Given a coding task, it can **explore the codebase, decide what to change, inspect results, make further changes until the task is completed.**


Examples: LLM Agents Across Application Areas

- **Science — Robin (FutureHouse):** multi-agent system supporting scientific research and discovery.
- **Law — Harvey Agents:** perform legal research, document analysis, due diligence, contract review and drafting.
- **Data & Business Analysis — Julius AI:** analyses and manipulates data, identifies patterns, and generates visualisations and reports.
- **Design & Prototyping — Figma Agent:** creates and edits designs and wireframes, uses design-system components, and turns static mockups into interactive prototypes through multi-step actions on the design canvas.

High-level LLM agent architecture



Concrete Agent Architecture



AI Agent Process / Loop

- **Perceive / Context:** Receive the user input and gather relevant context from the environment, including retrieved information and memory.
- **Reason / Plan:** Understand the task, consider what needs to be done, and decide the next step or sequence of actions.
- **Act:** Perform actions using tools, APIs, applications, other agents, or physical systems.
- **Observe:** Receive and interpret the results or changes produced by the actions.
- **Reflect / Correct:** Evaluate progress and, when needed, revise the plan or correct previous actions.
- **Learn / Update Memory:** Retain useful information from the interaction for subsequent steps or future interactions.

Model: The Foundation of an LLM Agent

- The **LLM / foundation model** provides the core **understanding, reasoning, and generation capabilities** used by the agent.
- Examples: **GPT, Claude, Gemini, Llama, Mistral, DeepSeek, Grok, ...**
- The model can:
 - **Understand and interpret** user requests and context
 - **Reason** about information and possible actions (work out what to do)
 - **Generate** text, code, queries, structured tool calls, and other artefacts
 - **Interpret** results / feedback returned by tools or the environment (e.g., requested data, success/failure, or an error message).
- The model itself does not necessarily have direct access to **tools, memory, or the external environment** — these capabilities are provided through the **agent architecture / harness**.

Example: How an Agent Uses the Model

- **User:** “Find my overdue invoices and prepare a reminder for each customer.”
- **Understand:** The LLM identifies that overdue invoices must be found and reminders prepared.
- **Reason / Decide:** The LLM determines that the invoice system needs to be queried.
- **Generate Tool Call:** The LLM generates/selects the appropriate tool call and parameters.
- **Execute Tool Call:** The agent runtime executes the API call and returns the results to the LLM.
- **Interpret:** The LLM interprets the returned invoice details (customers, amounts, due dates).
- **Generate Output:** The LLM drafts a reminder for each customer.

User Request → LLM → Tool Call → Agent Runtime/Tool → Tool Result → LLM → Output

Retrieval-Augmented Generation (RAG)

- **RAG provides the agent with relevant external information at task time.**
- **Information can come from:**
 - Documents and knowledge bases
 - Enterprise/private data
 - Web and other external sources
 - Search and retrieval systems
- **Retrieve:** Find information relevant to the current request.
- **Augment:** Add the retrieved information to the model's **context**.
- **Generate:** Use the retrieved information to generate a response or decide what to do next.

RAG enables agents to use **organisation-specific, private, specialised, or up-to-date information** that is not reliably available from the model alone.

Example: RAG in an Agent

The agent retrieves relevant information from external sources, adds it to the model's context, and uses it to generate a grounded response.



1

Retrieve

Find relevant information from internal

2

Augment

Add the retrieved information to the

3

Generate

The model uses this information to

Memory in LLM Agents

- **Memory enables an agent to retain and reuse information from previous interactions, actions, and outcomes.**
- **Short-term / working memory**
 - Keeps information needed while the agent is completing the current task or interaction
 - Allows the agent to use information from earlier steps when deciding or acting in later steps
- **Long-term / persistent memory**
 - Retains useful information **across interactions or sessions**
 - Can include user preferences, past decisions, experiences, and task-related information
- **Memory can be retrieved** when relevant and added to the agent's current context.
- **Memory can also be updated** as the agent obtains new useful information.
- **Why memory?**
 - Maintain continuity over time, avoid repeatedly asking for or processing the same information, use past experience and information to support future tasks.

Example: Using Memory in an Agent

Existing Memory (before this request)



User preferences

Prefers PDF format for reports.



Project information

FRAME EU project (P62518).
WP2 focus on benchmarking.



Past interactions

"We discussed the project progress and evaluation approach last week."



Other relevant information

...



User

"Prepare a report on the FRAME EU project. Use my preferred format. From now on, include a short executive summary at the beginning."

Retrieve from memory

Find relevant information (e.g., user preferences, project details, past interactions).

LLM Agent



Uses retrieved memory in the current context to generate a response.

Update memory

Store new useful information from this interaction, for example:

- new user preference
- new project information
- important follow-up

Updated Memory (after this request)



User preferences

Prefers PDF format for reports.
Include a short executive summary in future FRAME reports.



Project information

FRAME EU project (P62518).
WP2 focus on benchmarking.
Next review meeting on 20 September 2026.



Past interactions

"We discussed the project progress and evaluation approach last week."



Other relevant information

...




Agent

"I have prepared the report on the FRAME EU project (P62518) in PDF format. It includes a short executive summary at the beginning, as requested."



- **Tools enable an agent to interact with external systems and perform actions beyond model generation.**
- **Tools can include:**
 - APIs, databases, web/search services and applications
 - Code execution and command-line tools
 - Robots, IoT devices and other physical systems
- **Skills** provide reusable capabilities or procedures for accomplishing particular tasks, often combining **instructions, knowledge and/or tool use**.
- **MCP (Model Context Protocol)** provides a standard way for agents/AI applications to **connect to and use external tools, resources and services**.
- Through tools and skills, agents can:
 - **Retrieve information** from external systems
 - **Execute actions** and modify external state

Example: Using Tools and Skills via MCP



1

The agent requests a tool via MCP (e.g., `get_exchange_rate`).

2

The tool accesses the environment (exchange-rate service/API).

3

The environment returns the data (e.g., 1.08).

4

The tool returns the result to the agent, which provides the final answer to the user.

Context in LLM Agents

- **Context is the information made available to the model at a particular point during agent execution.**
- Context can include:
 - **User input** — prompts, text, voice, uploaded files
 - **System / agent instructions** — system prompt, role, constraints, policies
 - **Perception / environment information** — screenshots, images, audio, sensor data, application/system state
 - **Retrieved knowledge (RAG)** — relevant documents, data, or Web information
 - **Retrieved memory** — relevant information from previous interactions or tasks
 - **Tools / skills information** — available tools, descriptions, schemas, and instructions
 - **Observations / tool results** — API responses, execution results, errors, environment feedback
- The **agent/harness manages context**, selecting and assembling the information needed by the model for the current step.

Example: How Context Is Built for the LLM

- **User:** “Prepare a short update on the FRAME EU project for tomorrow’s meeting.”
- **Agent software / harness is configured to assemble:**
 - User request
 - Agent instructions
 - Relevant FRAME information from project documents
 - Relevant information from memory
 - Tomorrow’s meeting details from the calendar
 - Results from earlier task steps that are needed for the current step
- **This information is combined and provided to the LLM as context for the current step.**
- **The agent software / harness contains the logic that determines what information to retrieve and include in the LLM context.**

Planning / Task Decomposition in LLM Agents

- **Planning determines how an agent can accomplish a goal through a sequence of steps or actions.**
- For complex tasks, planning can involve:
 - **Task decomposition** — break a goal into smaller tasks or subgoals
 - **Action sequencing** — determine what actions should be performed and in what order
 - **Dependencies** — identify which steps depend on the results of others
 - **Tool/skill selection** — determine which capabilities may be needed for different steps
 - **Replanning** — update the plan when new information, results, or failures are observed
- **Not every task requires an explicit plan:** simple tasks may proceed directly to an action, while complex tasks may require multiple planned steps.
- Plans are typically **generated or updated using the LLM**, while the **agent/harness manages their execution and progress**

Example: Planning / Task Decomposition

Goal: “Help me prepare for my conference next week.”

Task decomposition — identify the tasks

- Arrange travel
- Prepare presentation
- Plan conference **attendance**

Planning — work out a sequence of actions to complete the tasks

- Travel: confirm conference location/dates, find flights, select flight, book
- Presentation: check presentation requirements, prepare slides, review slides
- Attendance: get conference programme, select sessions to attend, add them to calendar


Replanning: Adjust the plan if new information or problems arise during execution.

- **Software/control layer that coordinates the operations of the LLM agent.**
- **Orchestration & execution**
 - Coordinates the LLM, tools/skills, memory, RAG and other components
 - Controls execution of tasks and actions
- **Context & state management**
 - Builds and updates the context provided to the LLM
 - Tracks task state and progress
- **Monitoring & feedback**
 - Monitors actions, tool results, errors and execution progress
 - Coordinates reflection, external feedback and human review when needed
- **Safeguards / guardrails**
 - Check inputs, outputs and tool actions
 - Restrict actions based on rules and permissions

Example: Harness Managing an Agent

- **User:** “Find and book a flight to Paris for Monday morning.”
- **Context management**
 - Provides the LLM with the request plus relevant preferences:
depart from Dublin, after 07:00, budget €200
- **Orchestration & execution**
 - Runs the flight-search tool, returns the flight options to the LLM, then runs the booking tool after a flight is selected and confirmed.
- **Guardrails**
 - Enforces the rule: no purchase can be made without explicit user confirmation
- **Monitoring**
 - Records that the flight search succeeded, two options were returned, and the booking completed successfully.

LLM Agent Engineering Lifecycle



From Single-Agent to Multi-Agent Systems

- **Single LLM agent**

- Uses an LLM together with context, memory, tools/skills and a harness to perform tasks. One agent is responsible for accomplishing the task


- **Multi-agent system**

- Uses multiple LLM agents that interact and coordinate to accomplish a task or shared goal
- Agents may have different roles, capabilities, knowledge, tools or subgoals
- Work can be divided among agents, with different agents responsible for different parts of the task

- **Agent-to-Agent communication**

- Agents exchange information/results and coordinate their actions
- Agents need mechanisms to communicate, exchange information and delegate tasks
- Protocols such as A2A (Agent2Agent) support communication and collaboration between agents

Multi-Agent System Architecture



Example: Multi-Agent System for Software Development

- **User:** “Add a new feature to this application and make sure it works correctly.”
- **Coordinator Agent**
 - Decomposes the task and delegates work to specialised agents
- **Developer Agent**
 - Examines the codebase and implements the feature
- **Testing Agent**
 - Tests the implementation and identifies failures or problems
- **Review Agent**
 - Reviews the code and suggests corrections or improvements
- **Agents communicate and exchange results**, e.g. the Testing Agent sends detected failures back to the Developer Agent for correction.

Example: Multi-Agent System for Insurance Claims

- **User:** “Process this car insurance claim.”
- **Coordinator Agent**
 - Receives the claim and delegates work to specialised agents
- **Policy Agent**
 - Checks the customer’s policy, coverage and applicable conditions
- **Claims Assessment Agent**
 - Analyses the claim information and supporting documents
- **Fraud Detection Agent**
 - Checks the claim for potential fraud indicators
- **Agents communicate and exchange results**
 - The Coordinator Agent combines the findings and coordinates the next step...

Organising / Orchestrating Multiple Agents

- Multiple agents need a way to organise **who does what and how their work is coordinated**
- Common organisation patterns:
 - **Coordinator / Supervisor** — one agent coordinates the work and delegates tasks to other agents
 - **Peer-to-Peer** — agents communicate and collaborate directly with each other
 - **Hierarchical** — agents are organised at different levels, with some agents coordinating other agents
- The organisation determines **how tasks are distributed, how agents communicate, and how their work is coordinated.**

Example: Same Task, Different Multi-Agent Organisation

- **Task:** “Prepare a research report on AI agents.”
- **Coordinator / Supervisor**
 - Coordinator delegates search, analysis and writing to specialised agent
 - Agents return their results to the Coordinator, which combines them
- **Peer-to-Peer**
 - Search Agent passes papers to Analysis Agent
 - Analysis Agent passes findings directly to Writing Agent
- **Hierarchical**
 - Manager Agent coordinates Research and Writing lead agents
 - Research Lead Agent coordinates Search and Analysis agents
 - Writing Lead Agent produces the final report

- **Reliability & Quality**
 - Incorrect outputs, decisions, tool use or task execution
- **Safety & Security**
 - **Harmful/unintended actions** (e.g., deleting the wrong file), **prompt injection** (e.g., following malicious instructions in a webpage), and **interactions with untrusted or compromised tools** (e.g., a compromised tool asks the agent to send data to an external server).
- **Privacy & Confidentiality**
 - **Exposure or misuse of sensitive or confidential information through context, memory or tools**

- **Bias / Fairness**
 - Unfair or systematically different treatment of people/groups (e.g., based on gender or ethnicity)
- **Transparency**
 - It may be difficult to understand why the agent made a particular decision or took a particular action.
- **Autonomy & Control**
 - Agents can independently make decisions and take actions with real-world consequences
- **Safeguards & Quality Control**
 - **Testing & evaluation** (check behaviour), **guardrails & permissions** (control what the agent can do), **monitoring** (observe behaviour/actions), and **human oversight** (review/approve important actions)